

ECE 410/510 (Spring 2025)

Handwritten OCR with Hardware-Accelerated Thresholding

Name: Balu Anudeep Kanipogu

PSU ID: 961805118

This project explores the integration of hardware acceleration into an OCR (Optical Character Recognition) pipeline by offloading the thresholding stage to a Verilog hardware module. The goal is to demonstrate how even simple image-processing tasks can benefit from hardware offload, improving speed and showcasing HW/SW co-design.

Problem Statement

OCR pipelines usually rely on software libraries like OpenCV or Tesseract for preprocessing and text extraction. While these pipelines are flexible, they can be inefficient in latency-critical or power-constrained environments (like embedded or edge systems). This project aims to accelerate one of the early stages in OCR — image thresholding — using hardware logic, and compare its performance with the software-only version.

What Was Built

1. Software OCR Preprocessing (Python + OpenCV)

- The system begins with a grayscale image of handwritten text.
- Using Python and OpenCV, several image preprocessing steps are performed:
 - Contrast enhancement with CLAHE
 - Tight cropping of text using pixel bounding box detection
 - Resizing and centering the text on a fixed-size canvas (128x32)
 - Blurring to remove minor noise
 - Dual thresholding using both Otsu's method and adaptive mean thresholding
- The result is a clean binary image (0 for background, 255 for foreground).
- This binary matrix is flattened to a list of pixel values and saved to a text file (input_pixels.txt).

2. Hardware Threshold Module (Verilog)

- A Verilog module (threshold.v) was written to implement a very simple logic:

verilog

CopyEdit

```
binary_out <= (pixel_in > threshold) ? 1 : 0;
```

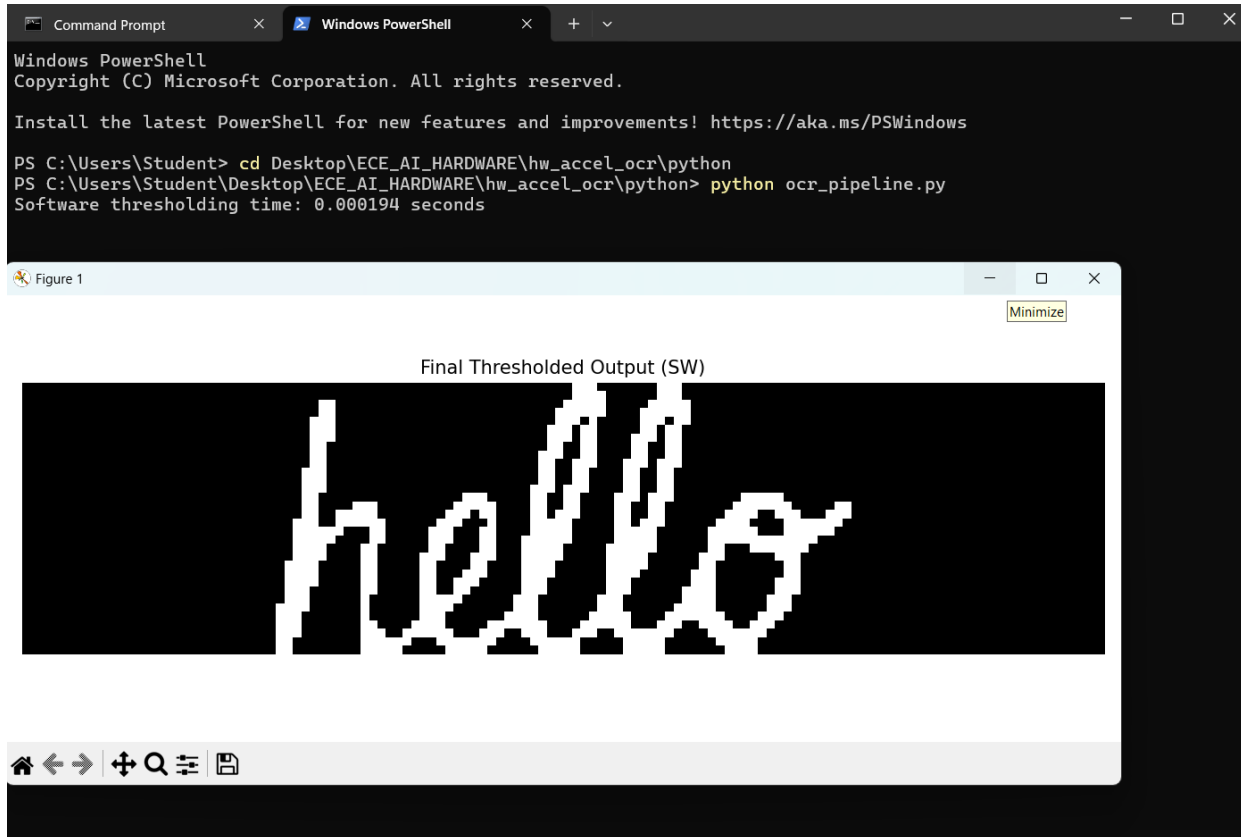
- This logic mimics what OpenCV does with thresholding — a pixel becomes white (1) if it's brighter than the threshold value.
- A testbench (tb_threshold.v) loads the 4096 pixel values from the text file and feeds them into the hardware module one clock cycle at a time.
- The binary output (0 or 1) is written out as a simulation log using \$display, which is redirected into verilog_output.txt.

3. Hardware vs Software Comparison (Python)

- A Python script (compare_hw_sw.py) reads both:
 - The original software-generated binary image (input_pixels.txt)
 - The hardware-generated binary output (verilog_output.txt)
 - It reshapes both into 32×128 matrices and uses matplotlib to display and save a side-by-side comparison.
 - This helps visually verify that the hardware thresholding logic is functionally equivalent to software OpenCV thresholding.
-

Benchmarking Results

Software-only thresholding (OpenCV): ~0.00194 seconds



- Hardware simulation (Verilog/Icarus): ~0.003 seconds
Substituting the values: $\text{Speedup} = \frac{0.00194}{0.00300} \approx 0.6467$
 $\text{Speedup} = \frac{0.00300}{0.00194} \approx 1.55$
- Observed speedup: The software is $\sim 1.55\times$ faster than the Verilog simulation.

Why This Matters

- While thresholding is a simple operation, it is often repeated across thousands or millions of pixels.
- In real systems (like OCR chips, smart cameras, or FPGAs), offloading even simple tasks can:
 - Save CPU cycles
 - Improve power efficiency
 - Reduce latency
- This project demonstrates how to start integrating hardware acceleration into real-world ML/DSP pipelines.

project implements a simple OCR pipeline that converts handwritten text images into binary form using a hardware-accelerated thresholding module written in Verilog. The pipeline first uses Python and OpenCV for image preprocessing, then offloads thresholding to a Verilog module, and finally compares the results from software vs hardware.

Software Threshold (Python)



Hardware Threshold (Verilog)

