

Find Bottlenecks (Analysis)

When you look at Q-learning, the major bottlenecks are:

Bottleneck Why

1. Q-table Lookup and Update Happens every action step; frequent memory access.
2. $\text{Max}(Q[\text{next_state}, :])$ calculation Need to scan all actions to find maximum reward — not parallelized.
3. Nested for-loops over episodes and steps Slow in Python because no GPU / parallelism.

Main bottleneck:

- 1) Finding max reward ($\text{np.max}(Q[\text{next_state}, :])$)
- 2) Updating Q-table entry ($Q[\text{state}, \text{action}]$)

These are the parts that dominate total runtime.

How to Propose a Hardware (HW) Acceleration

Idea:

Build a small hardware block that can parallel-scan the Q-table entries.

Hardware can fetch all actions for a state and find the maximum reward faster than software (Python loop).

Also update the Q-value instantly.

Hardware module will do:

Read all Q-values for next_state

Find $\text{max}(Q[\text{next_state}, :])$

Perform Q-value update (according to Bellman equation)

Simplify the logic: Think like a parallel max-finder + adder.

1. Bottleneck Analysis Bottleneck Analysis: In the Q-Learning FrozenLake implementation, the main computational bottlenecks are:

1: Q-table Lookup and Update: For each step, the agent accesses and updates the Q-value corresponding to the (state, action) pair, which involves memory read and write operations.

2:Max Operation: Calculating $\max(Q[\text{next_state}, :])$ requires scanning all possible actions for a given state, which is done sequentially in Python and is not parallelized.

3:Nested Loops: The outer loop over episodes and the inner loop over steps add to the computational load, but the max operation inside each step dominates.

Therefore, optimizing the $\max(Q[\text{next_state}, :])$ lookup and the Q-value update are key areas for acceleration.

2. Hardware Acceleration Proposal

Hardware Acceleration Proposal:

To speed up the Q-Learning process, we propose implementing a Q-table Accelerator Block in hardware. This block will:

- 1)Fetch all Q-values corresponding to the next state in parallel.
- 2)Perform a parallel max-finding operation across all actions.
- 3)Calculate the Q-value update using the Bellman update rule in hardware.
- 4)Write back the updated Q-value into the Q-table.

This will drastically reduce the time spent on the Q-table lookup and update, which are sequential and slow in Python. By using comparators, adders, and small memory blocks in hardware, the critical path is shortened significantly.